

Zero-Trust Architecture for MCP-Based AI Agents: Continuous Identity and Behavioral Monitoring

Yoshihiro Ando
Independent Researcher
yosh5295@gmail.com
Tokyo, Japan

Abstract—As autonomous AI agents evolve into distributed system operators via the Model Context Protocol (MCP), the assumption of network-level security becomes obsolete. This paper, representing Phase 2: Zero Trust (Execution Security) of a comprehensive security trilogy, presents a framework for securing agentic workloads through "Spatial Security" and continuous behavioral monitoring. Building upon the structural guardrails of Phase 1, we implement a Zero Trust Architecture (ZTA) centered on asymmetric workload identity (RSA) and a Mamba-based Reasoning Sentinel. The sentinel utilizes a NumPy-only State Space Model (SSM) to perform real-time anomaly detection on the agent's reasoning traces, providing a behavioral Intrusion Detection System (IDS) for AI logic. Our results demonstrate that the combination of sub-millisecond cryptographic verification and lightweight behavioral monitoring effectively mitigates risks like prompt injection and semantic drift without impeding agentic responsiveness.

Index Terms—Zero Trust Architecture, AI Agents, Model Context Protocol, Workload Identity, Mamba Sentinel, Intrusion Detection System (IDS), Reasoning Integrity.

I. INTRODUCTION

Phase 1 of this trilogy [2] established structural containment via sandboxing and "No-Shell" architectures. However, a "jailbroken" agent might still perform harmful actions that are syntactically valid but semantically malicious. For instance, an agent permitted to "manage documents" might be tricked into deleting a critical database file if it can be represented as a document. This necessitates a shift from static containment to dynamic, behavioral verification.

We propose a Zero Trust Architecture (ZTA) for AI agents, where the principle of "Never Trust, Always Verify" is applied not just to the network, but to the *intent* and *identity* of the AI's reasoning process.

This paper is the second installment of our security trilogy:

- **Phase 1: Guardrails (Pre-execution Security)** [2]. Structural containment and sandboxing.
- **Phase 2: Zero Trust (Execution Security)**. The current work, focusing on **Continuous Verification** via identity propagation and behavioral integrity.
- **Phase 3: Post-Quantum (Temporal Security)** [3]. Long-term integrity and quantum-resistant governance.

II. ASYMMETRIC IDENTITY AND TOOL ORCHESTRATION

In a Zero Trust model, an agent's authority must be tied to a verifiable identity rather than its session or location.

A. Workload Identity (RSA-2048)

`llm-cli` assigns each agent instance a unique RSA-2048 key pair. When the agent requests a tool call via MCP, the system generates a signed **Identity Token (JWT)**. This token acts as a "Passport" for the agent, containing:

- **sub (Subject)**: The agent's unique ID (e.g., `user@hostname`).
- **roles**: Assigned RBAC roles (e.g., `admin`, `user`).
- **iat/exp**: Issued-at and expiration timestamps to prevent replay attacks.

This ensures that remote MCP servers can cryptographically verify that the tool call originated from a legitimate, untampered agent instance.

B. Role-Based Access Control (RBAC)

Permissions are decoupled from the LLM's prompt. Even if the LLM "believes" it has root access due to a prompt injection, the underlying security engine enforces a strict RBAC policy. If a tool (e.g., `edit_file`) requires the `admin` role, the execution is blocked regardless of the AI's "monologue" unless the signed identity token carries that claim.

III. BEHAVIORAL IDS: THE MAMBA SENTINEL

While identity ensures **who** is calling the tool, it does not ensure **why**. To detect semantic attacks like prompt injection, we introduce the **Mamba Sentinel**.

A. The Need for Lightweight Monitoring

Traditional monitoring using secondary LLMs (e.g., using GPT-4o-mini to verify GPT-4o's output) is slow (1-2 seconds of latency) and expensive. We utilize a **Mamba (State Space Model)** implementation [1] optimized with pure NumPy. Unlike Transformers, Mamba's linear scaling and recurrent state make it ideal for monitoring streaming reasoning traces at a byte level with sub-millisecond overhead.

B. Reasoning Integrity and Surprise Score

The Sentinel monitors the LLM's "internal monologue" (reasoning tokens) in real-time. It calculates a **Surprise Score** based on the cross-entropy loss between the Sentinel's predicted next-token distribution and the agent's actual output.

$$L_{\text{sentinel}} = - \sum P_{\text{sentinel}}(x_t | h_{t-1}) \log P_{\text{agent}}(x_t) \quad (1)$$

A sudden spike in "surprise" indicates a statistical anomaly or "Semantic Drift." This often happens when the agent transitions from a helpful assistant state to an adversarial state (e.g., suddenly generating obfuscated code or trying to access sensitive files after a prompt injection).

C. Dynamic Intervention (Green/Yellow/Red)

The system maintains a real-time status:

- **Green:** Normal operation.
- **Yellow:** Minor anomaly detected. The system increases the entropy-scanning sensitivity and alerts the user.
- **Red:** Significant anomaly. The system **automatically suspends execution** and forces a Human-in-the-Loop (HITL) review.




Fig. 1. Phase 2 Architecture: Continuous verification via Asymmetric Identity and Mamba Behavioral Sentinel.

IV. EVALUATION

We evaluated the Zero Trust framework on an **AMD Ryzen 5 8540U** system.

A. Execution Latency (Empirical)

Table I summarizes the latencies. The Mamba Sentinel adds only 10ms to the reasoning process, which is nearly two orders of magnitude faster than a "Dual-LLM" verifier approach.

TABLE I
ZERO TRUST AND IDS EXECUTION LATENCY

Component	Latency (ms)	Type
Identity Verification (RS256)	0.0296	Cryptographic
Mamba Sentinel (SSM)	10.3122	ML-Inference

B. Detection Accuracy

In testing with known prompt injection techniques (e.g., "DAN", "Ignore previous instructions"), the Mamba Sentinel detected escapes by identifying statistical anomalies in the reasoning trace before a harmful tool call was even attempted.

This provides a "Defense in Depth" that complements the structural guardrails of Phase 1.

V. CONCLUSION

By decoupling reasoning from authority through asymmetric identities and providing real-time behavioral monitoring via the Mamba Sentinel, we ensure that autonomous agents remain "Verified by Identity and Monitored by Logic." This fulfills the promise of agentic productivity with robust, low-latency execution security. The final phase of our trilogy addresses resilience against future quantum threats and remote integrity proofs.

REFERENCES

- [1] A. Gu and T. Dao, "Mamba: Linear-Time Sequence Modeling with Selective State Spaces," 2023.
- [2] Y. Ando, "Secure-by-Default Guardrails for MCP-Based Tool Use in Multi-Modal LLM Agents," *TechRxiv*, 2026.
- [3] Y. Ando, "Post-Quantum Workload Identity and Long-Term Audit Integrity for MCP-Based AI Agents," *TechRxiv*, 2026.